CMPE 493
INTRODUCTION TO INFORMATION RETRIEVAL

Introduction

Arzucan Özgür

Department of Computer Engineering, Boğaziçi University
September 29, 2015

Course Staff

- Instructor: Arzucan Özgür
  - Office: ETA 18
  - Phone: 0212-359-7226
  - E-mail: arzucan.ozgur@boun.edu.tr

  (Please include CMPE493 in your subject when sending e-mail.)
  - Office hours: Monday 14:00-15:00, Tuesday 13:00-15:00, or by appointment.
- TAs:
- Şaziye Betül Özateş (sbetulbilgin@gmail.com)
- Alper Çetiner (alper.cetiner@boun.edu.tr)

Text book

- Christopher D. Manning, Prabhakar Raghavan and Hinrich Schütze, Introduction to Information Retrieval, Cambridge University Press. 2008.

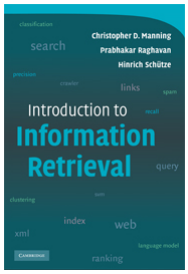

Available online (free) at the website of the book:
http://nlp.stanford.edu/IR-book/information-retrieval-book.html

Reference book (Optional)

- Daniel Jurafsky and James H. Martin, SPEECH and LANGUAGE PROCESSING: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition, Second Edition, 2008.

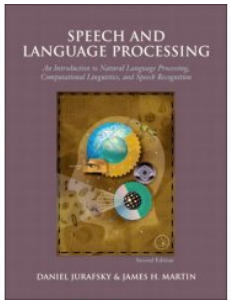

Available at the Bookstore.

Reference book (Optional)

- Christopher D. Manning and Hinrich Schütze, Foundations of Statistical Natural Language Processing, MIT Press, 1999. http://nlp.stanford.edu/fsnlp/

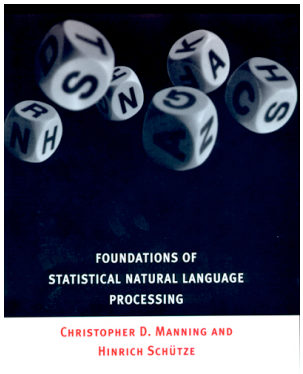

Course Web Site:

- We will use the Moodle Course Management System for lecture notes, announcements, homework/project submissions, and grading.
  - https://moodle.boun.edu.tr

You will automatically be subscribed to the system. You can login using your "boun" e-mail account's username and password.

Grading

- Midterm Exam: 15%
- Final Exam: 15%
- 3-4 Assignments: 30%
- Term Project: 35%
- Class Participation: 5%

Grading – Exams

- In-class midterm and final exams
- Consisting of problems covering the lecture material
- Closed book/notes
- Dates:
  - Midterm Exam: November 4, in the lecture hour (15:00-17:00)
  - Final Exam: As scheduled by the registration office

Grading – Homework Assignments

- Involve some programming where you will implement and test some of the techniques that we cover in class.
- You can use any programming language of your choice such as Perl, Python, Java, etc.
- We should be able to run your program.
- You should provide a readme file, explaining how to run your program.

Term Project

One of the aspects of this course is preparing you for original research in IR.

- Identifying an interesting problem
- Gathering relevant literature and datasets
- Solving it using new algorithms
- Evaluating the results

- Ability to present your ideas and research
  - Writing up your results in a scientific paper format
  - Presenting a research talk to a scientific audience

Term Project

- The project teams can consist of one or two people. (Teams consisting of two people is recommended)
- Each team will choose a project topic by selecting a recent scientific paper from an IR/NLP conference or journal.
- The project will involve replicating the work done in the paper and proposing extensions/improvements to the existing work. The proposed extensions do not need to be implemented.

Some of the Relevant Scientific Conferences

- ACM SIGIR Conference on Research and Development in Information Retrieval
- Conference on Information and Knowledge Management (CIKM)
- ACM International Conference on Web Search and Web Data Mining (WSDM)
- Association for Computational Linguistics (ACL)
- North American Association for Computational Linguistics (NAACL)
- Empirical Methods in Natural Language Processing (EMNLP)
- International Conference on Computational Linguistics (COLING)
- You can select your papers from relevant journals as well, including Information Retrieval, Computational Linguistics, TACL, Natural Language Engineering, and Journal of the Association for Information Science and Technology (JASIST)

Term Project - Deliverables

- Paper selection and 1-2 page description of the methodology planned to be used to replicate the work
  - Date: November 2
- Short project presentation in the end of the semester
  - Tentative Dates: December 15-16, December 22-23, lecture hours
- Submit a project report in the end of the semester.
  - Tentative Date: Final exam date

Information Retrieval

- Information Retrieval (IR) is finding material (usually documents) of an unstructured nature (usually text) that satisfies an information need from within large collections.

Basic assumptions of Information Retrieval

- Collection: Fixed set of documents
- Goal: Retrieve documents with information that is relevant to the user's information need and helps the user complete a task

The classic search model

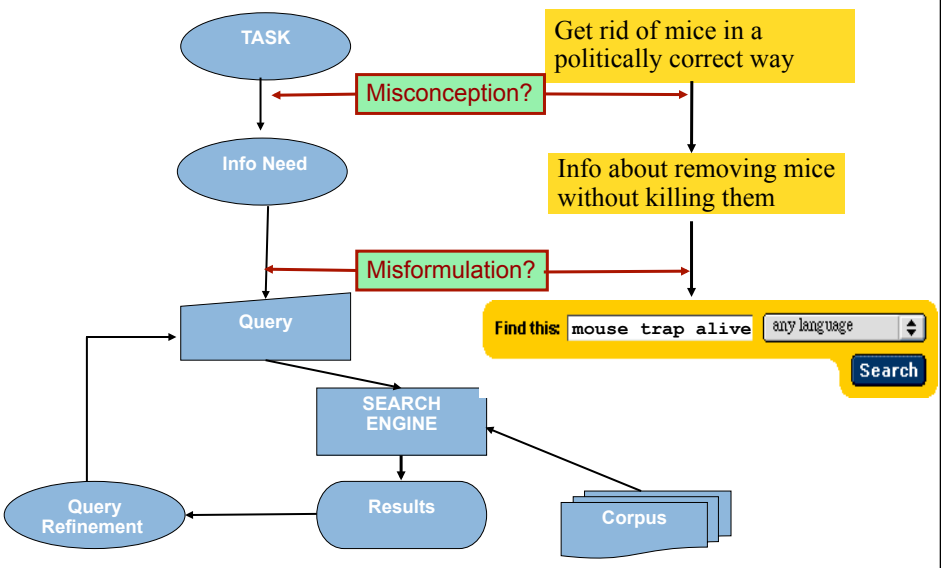

How good are the retrieved docs?

- *Precision*: Fraction of retrieved docs that are relevant to user's information need
- *Recall*: Fraction of relevant docs in collection that are retrieved

- More precise definitions and measurements to follow in later lectures

Examples of search engines

- Conventional (library catalog).
  Search by keyword, title, author, etc.
- Text-based (Google, Yahoo!, Bing, Yandex, Baidu; also email search, laptop search etc.)
  Search by keywords. Limited search using queries in natural language.
- Multimedia (QBIC, WebSeek)
  Search by visual appearance (shapes, colors,…).
- Question answering systems (Ask, NSIR, Answerbus)
  Search in (restricted) natural language
- Other:
  music retrieval

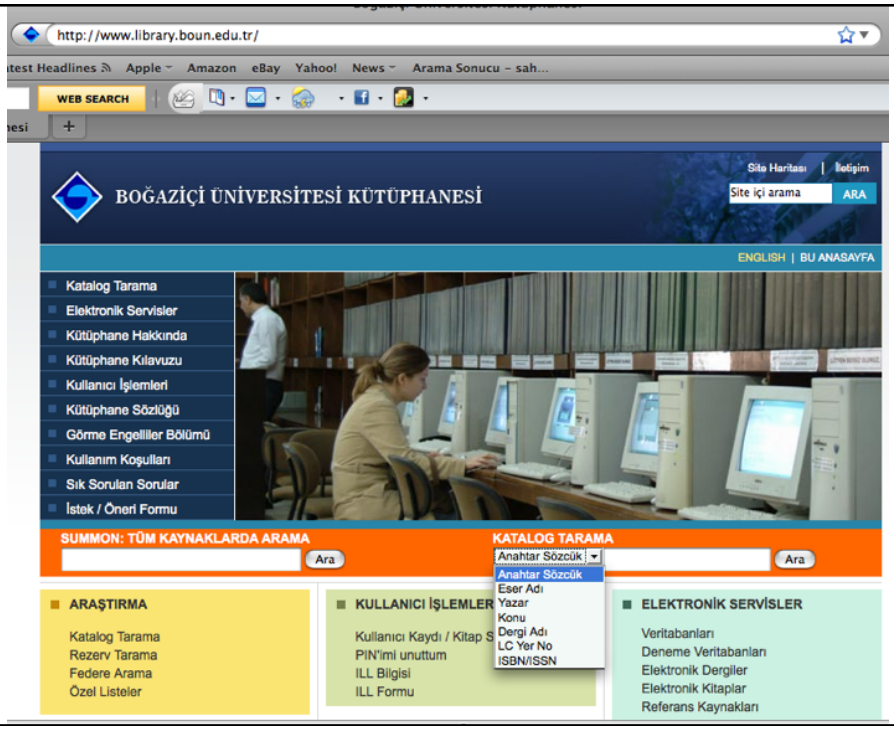

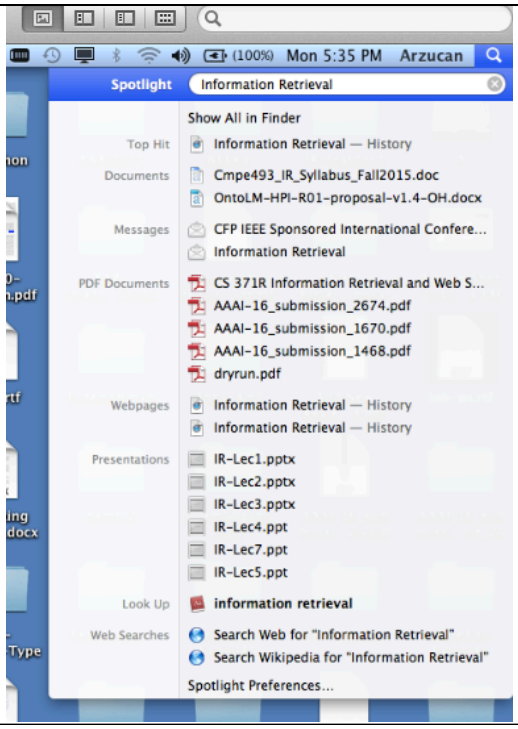

Laptop search
e.g. Spotlight

IR systems on the Web

- Search for Web pages: http://www.google.com
- Domain specific search (e.g., legal, biomedical): PubMed
- Search for images: http://www.picsearch.com
- Search for image content: http://wang14.ist.psu.edu/
- Search for answers to questions: http://www.ask.com
- Music retrieval: http://www.rotorbrain.com/foote/musicr/

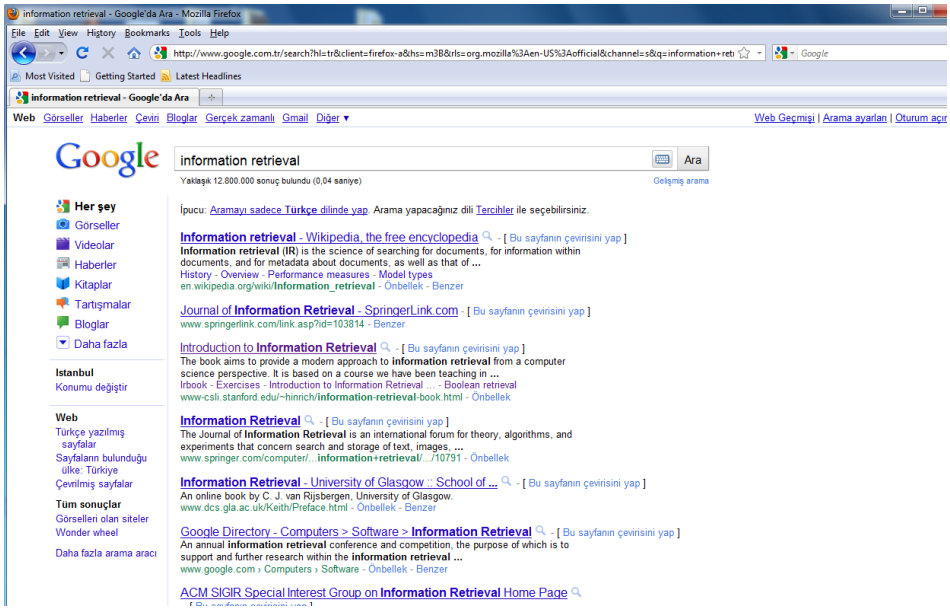

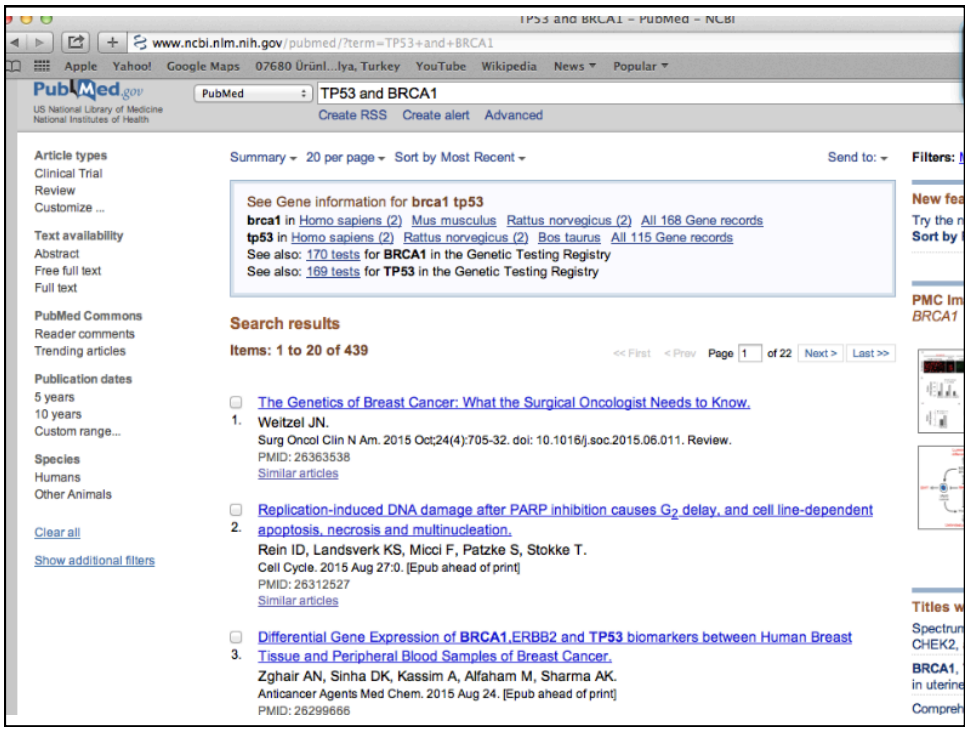
TP53 and BRCA1 - PubMed - NCBI
www.ncbi.nlm.nih.gov/pubmed/?term=TP53+and+BRCA1
Apple Yahoo! Google Maps 07680 Ürünl...lya, Turkey YouTube Wikipedia News Popular
PubMed.gov
US National Library of Medicine National Institutes of Health
PubMed
TP53 and BRCA1
Create RSS Create alert Advanced
Article types
Clinical Trial
Review
Customize ...
Text availability
Abstract
Free full text
Full text
PubMed Commons
Reader comments
Trending articles
Publication dates
5 years
10 years
Custom range...
Species
Humans
Other Animals
Clear all
Show additional filters
Summary 20 per page Sort by Most Recent
Send to:
Filters:
See Gene information for brca1 tp53
brca1 in Homo sapiens (2) Mus musculus Rattus norvegicus (2) All 168 Gene records
tp53 in Homo sapiens (2) Rattus norvegicus (2) Bos taurus All 115 Gene records
See also: 170 tests for BRCA1 in the Genetic Testing Registry
See also: 169 tests for TP53 in the Genetic Testing Registry
Search results
Items: 1 to 20 of 439
<< First < Prev Page 1 of 22 Next > Last >>
1. The Genetics of Breast Cancer: What the Surgical Oncologist Needs to Know.
Weitzel JN.
Surg Oncol Clin N Am. 2015 Oct;24(4):705-32. doi: 10.1016/j.soc.2015.06.011. Review.
PMID: 26363538
Similar articles
2. Replication-induced DNA damage after PARP inhibition causes G2 delay, and cell line-dependent apoptosis, necrosis and multinucleation.
Rein ID, Landsverk KS, Micci F, Patzke S, Stokke T.
Cell Cycle. 2015 Aug 27:0. [Epub ahead of print]
PMID: 26312527
Similar articles
3. Differential Gene Expression of BRCA1,ERBB2 and TP53 biomarkers between Human Breast Tissue and Peripheral Blood Samples of Breast Cancer.
Zghair AN, Sinha DK, Kassim A, Alfaham M, Sharma AK.
Anticancer Agents Med Chem. 2015 Aug 24. [Epub ahead of print]
PMID: 26299666

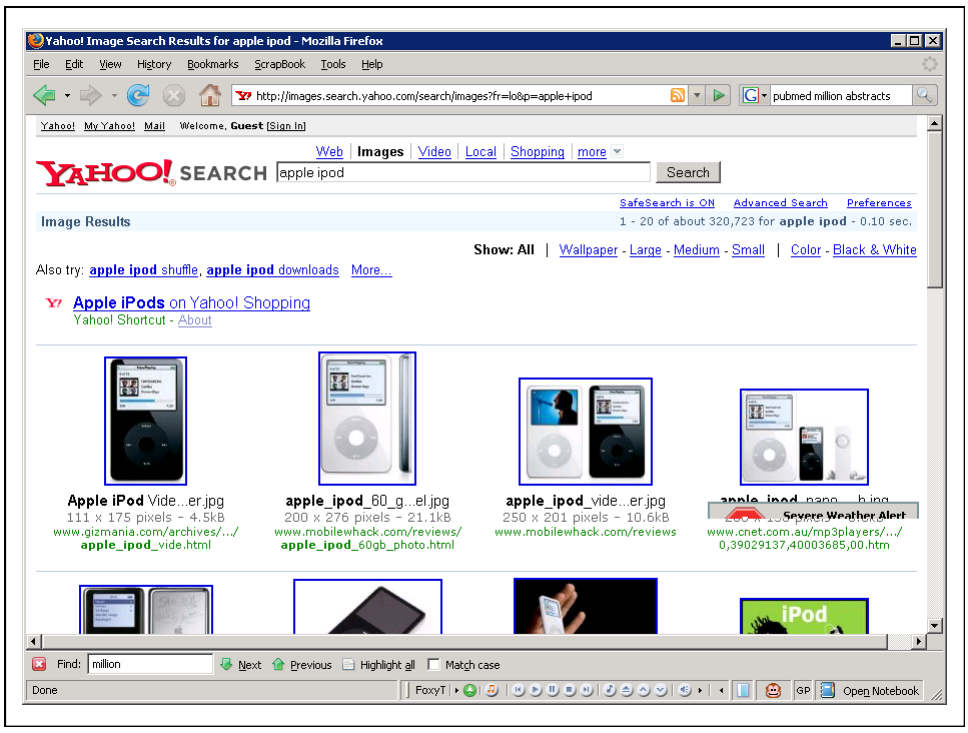
Yahoo! Image Search Results for apple ipod - Mozilla Firefox
File Edit View History Bookmarks ScrapBook Tools Help
http://images.search.yahoo.com/search/images?fr=lo&p=apple+ipod
pubmed million abstracts
Yahoo! My Yahoo! Mail Welcome, Guest [Sign In]
Web | Images | Video | Local | Shopping | more
YAHOO! SEARCH apple ipod Search
SafeSearch is ON Advanced Search Preferences
Image Results
1 - 20 of about 320,723 for apple ipod - 0.10 sec.
Show: All | Wallpaper - Large - Medium - Small | Color - Black & White
Also try: apple ipod shuffle, apple ipod downloads More...
Apple iPods on Yahoo! Shopping
Yahoo! Shortcut - About
Apple iPod Vide...er.jpg
111 x 175 pixels - 4.5kB
www.gizmania.com/archives/.../apple_ipod_vide.html
apple_ipod_60_g...el.jpg
200 x 276 pixels - 21.1kB
www.mobilewhack.com/reviews/apple_ipod_60gb_photo.html
apple_ipod_vide...er.jpg
250 x 201 pixels - 10.6kB
www.mobilewhack.com/reviews
www.cnet.com.au/mp3players/.../0,39029137,40003685,00.htm
Severe Weather Alert
iPod
Find: million Next Previous Highlight all Match case
Done
FoxyT GP Open Notebook

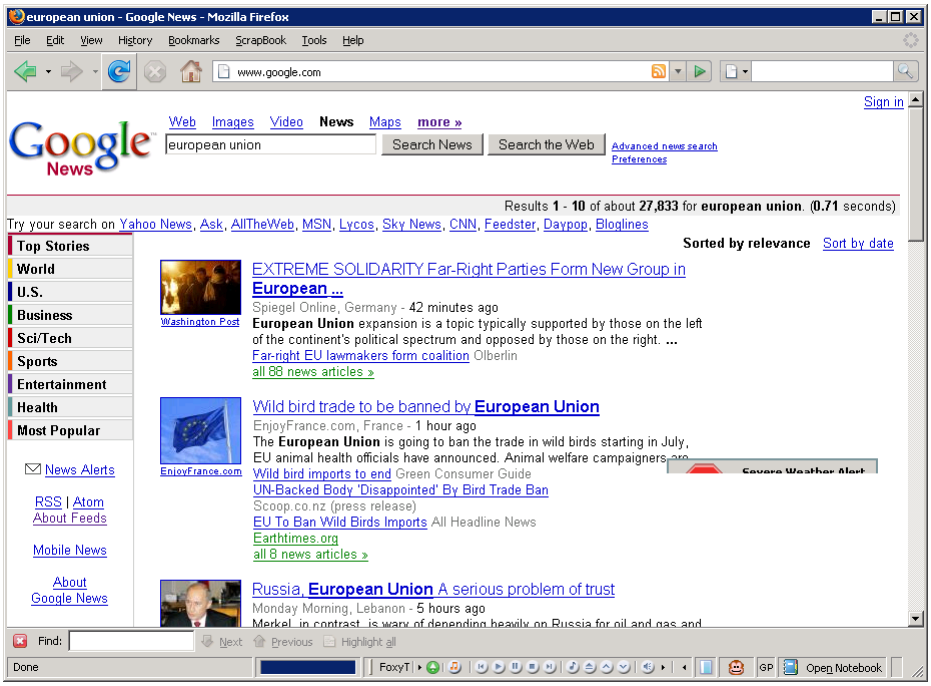
european union - Google News - Mozilla Firefox
File Edit View History Bookmarks ScrapBook Tools Help
www.google.com
Sign in
Google News
Web Images Video News Maps more »
european union
Search News
Search the Web
Advanced news search
Preferences
Results 1 - 10 of about 27,833 for european union. (0.71 seconds)
Try your search on Yahoo News, Ask, AllTheWeb, MSN, Lycos, Sky News, CNN, Feedster, Daypop, Bloglines
Top Stories
World
U.S.
Business
Sci/Tech
Sports
Entertainment
Health
Most Popular
News Alerts
RSS | Atom
About Feeds
Mobile News
About Google News
Sorted by relevance
Sort by date
EXTREME SOLIDARITY Far-Right Parties Form New Group in European ...
Washington Post
Spiegel Online, Germany - 42 minutes ago
European Union expansion is a topic typically supported by those on the left of the continent's political spectrum and opposed by those on the right. ...
Far-right EU lawmakers form coalition Olberlin
all 88 news articles »
Wild bird trade to be banned by European Union
EnjoyFrance.com
EnjoyFrance.com, France - 1 hour ago
The European Union is going to ban the trade in wild birds starting in July,
EU animal health officials have announced. Animal welfare campaigners are
Wild bird imports to end Green Consumer Guide
UN-Backed Body 'Disappointed' By Bird Trade Ban
Scoop.co.nz (press release)
EU To Ban Wild Birds Imports All Headline News
Earthtimes.org
all 8 news articles »
Russia, European Union A serious problem of trust
Monday Morning, Lebanon - 5 hours ago
Find: Next Previous Highlight all
Done
FoxyT
GP
Open Notebook

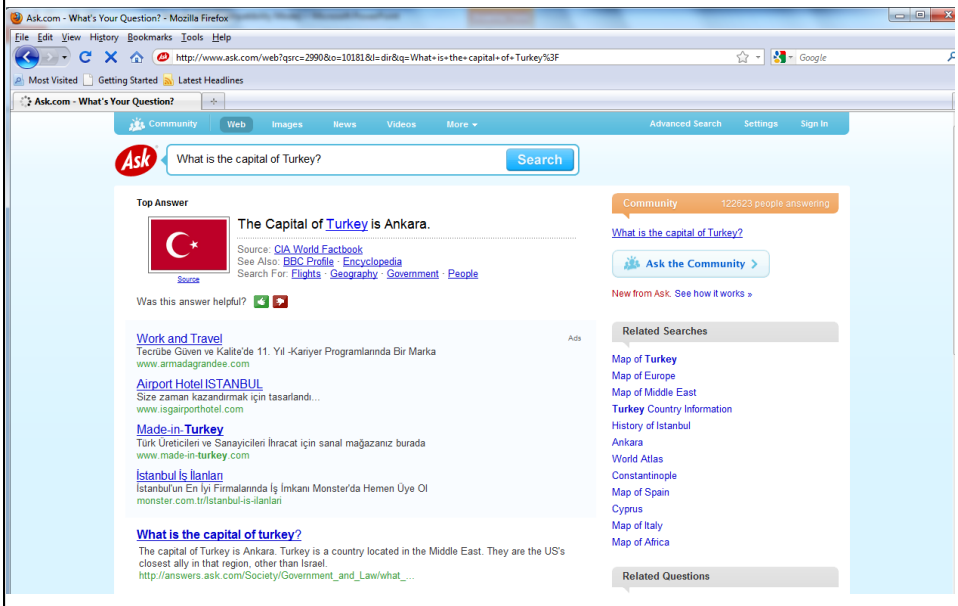
Ask.com - What's Your Question? - Mozilla Firefox
File Edit View History Bookmarks Tools Help
http://www.ask.com/web?qsrc=2990&o=10181&l=dir&q=What+is+the+capital+of+Turkey%3F
Most Visited Getting Started Latest Headlines
Ask.com - What's Your Question?
Community Web Images News Videos More
Advanced Search Settings Sign In
Ask
What is the capital of Turkey?
Search
Top Answer
The Capital of Turkey is Ankara.
Source: CIA World Factbook
See Also: BBC Profile · Encyclopedia
Search For: Flights · Geography · Government · People
Source
Was this answer helpful?
Work and Travel
Tecrübe Güven ve Kalite'de 11. Yıl -Kariyer Programlarında Bir Marka
www.armadagrandee.com
Airport Hotel ISTANBUL
Size zaman kazandırmak için tasarlandı...
www.isgairporthotel.com
Made-in-Turkey
Türk Üreticileri ve Sanayicileri İhracat için sanal mağazanız burada
www.made-in-turkey.com
İstanbul İş İlanları
İstanbul'un En İyi Firmalarında İş İmkanı Monster'da Hemen Üye Ol
monster.com.tr/istanbul-is-ilanlari
Ads
What is the capital of turkey?
The capital of Turkey is Ankara. Turkey is a country located in the Middle East. They are the US's closest ally in that region, other than Israel.
http://answers.ask.com/Society/Government_and_Law/what_...
Community 122623 people answering
What is the capital of Turkey?
Ask the Community
New from Ask. See how it works »
Related Searches
Map of Turkey
Map of Europe
Map of Middle East
Turkey Country Information
History of Istanbul
Ankara
World Atlas
Constantinople
Map of Spain
Cyprus
Map of Italy
Map of Africa
Related Questions

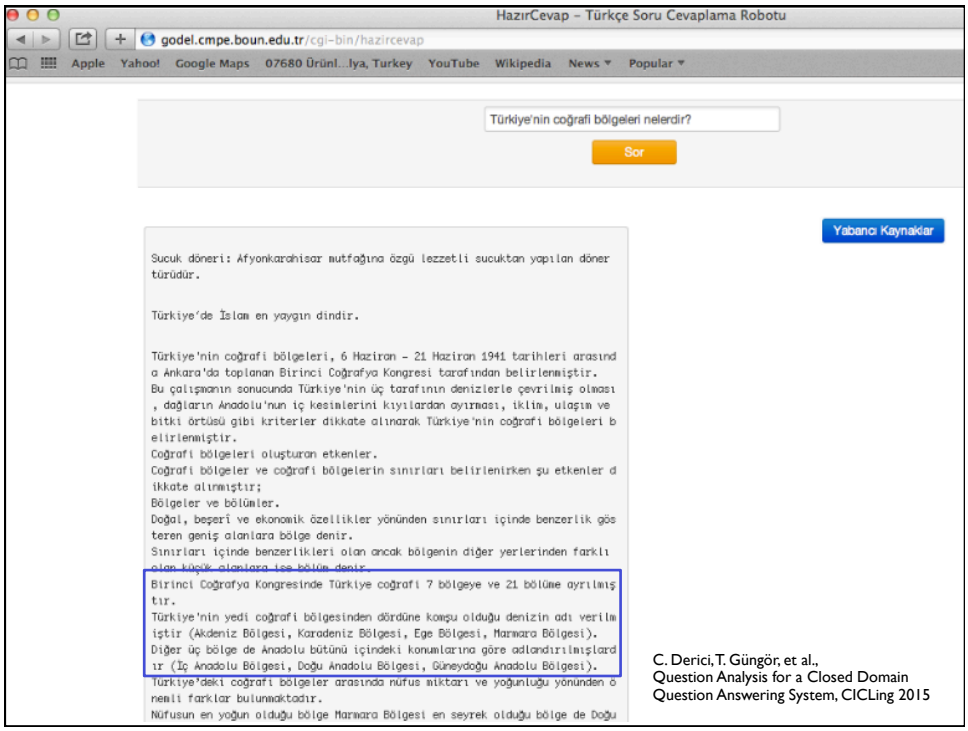

C. Derici, T. Güngör, et al.,
Question Analysis for a Closed Domain
Question Answering System, CICLing 2015

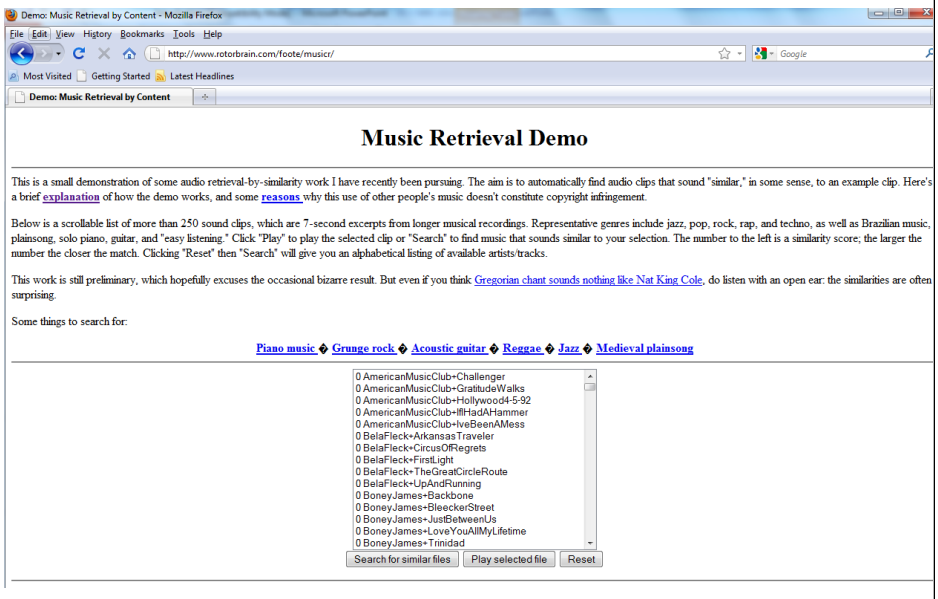

What does it take to build a search engine?

- Decide what to index
- Collect it
- Index it (efficiently)
- Keep the index up to date
- Provide user-friendly query facilities

What else?

- Understand the structure of the web for efficient crawling
- Understand user information needs
- Preprocess unstructured textual data
- Cluster data
- Classify data
- Evaluate performance

Goals of the course

- Understand how search engines work
- Understand the limits of existing search technology
- Learn about the state of the art in IR research
- Learn to analyze textual data sets
- Learn to evaluate information retrieval systems
- Learn about standardized document collections
- Learn about text similarity measures
- Learn about semantic dimensionality reduction
- Learn about web crawling
- Learn to use existing software
- Understand the dynamics of the Web by building appropriate mathematical models
- Build working systems that assist users in finding useful information from large collections

Topics (tentative list)

- Boolean model; text pre-processing; inverted indexes
- Approximate string matching and tolerant retrieval
- Index construction and compression
- Vector space model; text-similarity metrics; term weighting; ranked retrieval
- Evaluating information retrieval systems
- Relevance feedback; query expansion
- Language models for information retrieval
- Text classification and clustering
- Latent semantic indexing
- Web search and crawling
- Link analysis (e.g. hubs and authorities, Google PageRank)

References

- Content adapted from Prof. Dragomir Radev and the IR book's web site.