
CMPE 493
INTRODUCTION TO
INFORMATION RETRIEVAL

Phrase Queries and
Positional Indexes

Department of Computer Engineering, Boğaziçi University
September 30, 2015

Phrase queries

- ▶ Want to be able to answer queries such as “**Boğaziçi University**” – as a phrase
- ▶ Thus the sentence “*Sabancı University is at the other side of Boğaziçi Bridge.*” is not a match.
- ▶ For this, it no longer suffices to store only
<term : docs> entries

A first attempt: Biword indexes

- ▶ Index every consecutive pair of terms in the text as a phrase
- ▶ For example the text “Friends, Romans, Countrymen” would generate the biwords
 - ▶ *friends romans*
 - ▶ *romans countrymen*
- ▶ Each of these biwords is now a dictionary term
- ▶ Two-word phrase query-processing is now immediate.



Longer phrase queries

- ▶ ***boğaziçi university students*** can be broken into the Boolean query on biwords:
boğaziçi university AND university students

Without the docs, we cannot verify that the docs matching the above Boolean query do contain the phrase.

Can have false positives!



Extended biwords

- ▶ Parse the indexed text and perform part-of-speech (POS) tagging.
- ▶ Bucket the terms into (say) Nouns (N) and articles/prepositions (X).
- ▶ Call any string of terms of the form NX*N an extended biword.
 - ▶ Each such extended biword is now made a term in the dictionary.
- ▶ Example: **catcher in the rye**

N X X N
- ▶ Query processing: parse it into N's and X's
 - ▶ Segment query into enhanced biwords
 - ▶ Look up in index: **catcher rye**

Issues for biword indexes

- ▶ False positives, as noted before
- ▶ Index blowup due to bigger dictionary
 - ▶ Infeasible for more than biwords, big even for them
- ▶ Biword indexes are not the standard solution (for all biwords) but can be part of a compound strategy

Solution 2: Positional indexes

- ▶ In the postings, store, for each **term** the position(s) in which tokens of it appear:

<**term**, number of docs containing **term**;

doc1: position1, position2 ... ;

doc2: position1, position2 ... ;

etc.>



Positional index example

<**be**: 993427;

1: 7, 18, 33, 72, 86, 231;

2: 3, 149;

4: 17, 191, 291, 430, 434;

5: 363, 367, ...>



Which of docs **1,2,4,5**
could contain "*to be*
or not to be"?

- ▶ For phrase queries, we use a merge algorithm recursively at the document level
- ▶ But we now need to deal with more than just equality



Processing a phrase query

- ▶ Extract inverted index entries for each distinct term: **to**, **be**, **or**, **not**.
- ▶ Merge their *doc:position* lists to enumerate all positions with “**to be or not to be**”.
 - ▶ **to**:
 - ▶ 2:1,17,74,222,551; 4:8,16,190,429,433; 7:13,23,191; ...
 - ▶ **be**:
 - ▶ 1:17,19; 4:17,191,291,430,434; 5:14,19,101; ...
- ▶ Same general method for proximity searches

Proximity queries

- ▶ LIMIT! /3 STATUTE /3 FEDERAL /2 TORT
 - ▶ Again, here, /k means “within k words of”.
- ▶ Clearly, positional indexes can be used for such queries; biword indexes cannot.

Positional index size

- ▶ You can compress position values/offsets: we'll talk about that later.
- ▶ Nevertheless, a positional index expands postings storage *substantially*
- ▶ Nevertheless, a positional index is now standardly used because of the power and usefulness of phrase and proximity queries ... whether used explicitly or implicitly in a ranking retrieval system.



Rules of thumb

- ▶ A positional index is 2–4 as large as a non-positional index
- ▶ Positional index size 35–50% of volume of original text
- ▶ Caveat: all of this holds for “English-like” languages



Combination schemes

- ▶ These two approaches can be profitably combined
 - ▶ For particular phrases (“**Michael Jackson**”, “**Britney Spears**”) it is inefficient to keep on merging positional postings lists
 - ▶ Even more so for phrases like “**The Who**”



References

- ▶ *Introduction to Information Retrieval*, chapter 2
 - ▶ <http://nlp.stanford.edu/IR-book/information-retrieval-book.html>

